

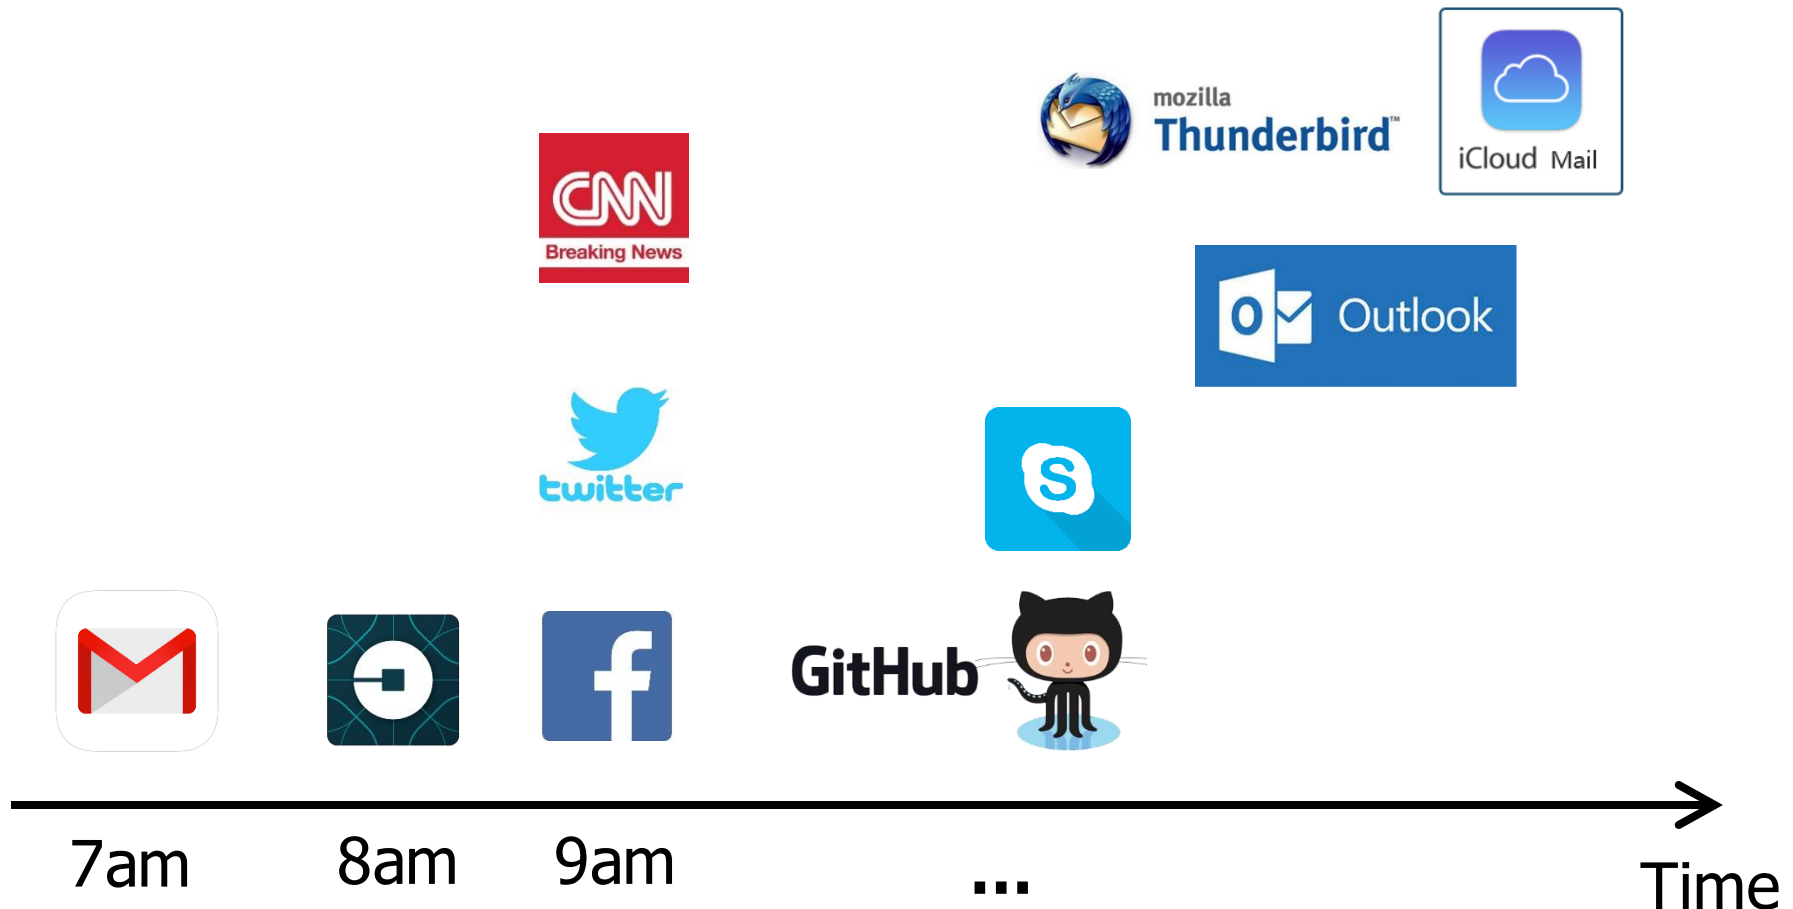
Programmable In-network Security

A vision for network security in the next generation

Ang Chen

University of Michigan

The modern life runs on network services



- We rely on network services every day!

However, networks are plagued with attacks



'Internet Of Things' hacking a
wide

SCIENTIFIC
AMERICAN

Power grid cyber attacks keep the
Pentagon up at night



AWS outage knocks Amazon, Netflix,
Tinder and IMDb in MEGA data c

The New York Times

Cyberattacks: Stuxnet and Flame

Forbes

Amazon.com goes
\$66,240 per minute



Global DNS outage
hits Microsoft Azure customers

owns cost countries \$2.4



Hacked home devices caused
massive Internet outage

billion last year

Skype calling knocked online by



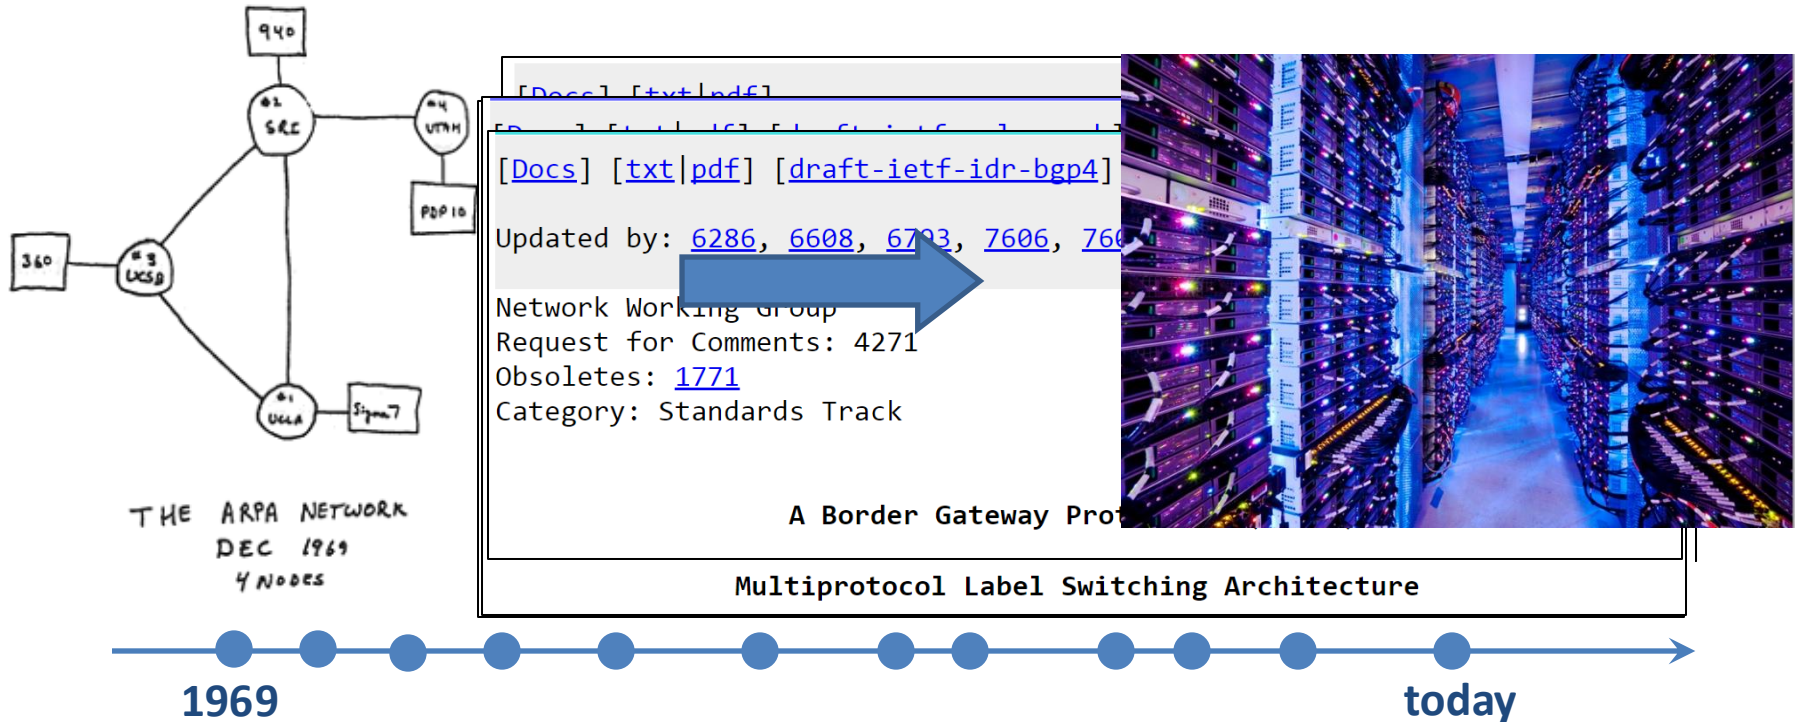
Azure Nightmare: Cust
Intermittent Outage

BUSINESS INSIDER

'Lizard Squad' took down Xbox Live
and PlayStation network

'Maintenance' caused
nationwide internet outage

How did we get here?




- Networks have evolved significantly over the years

Challenge: network security




- Our ability to secure them hasn't caught up.
 - The network infrastructure itself doesn't have security support



"The Internet is brittle and fragile and too easy to take down. It's a conduit for criminal activity."

--- Vint Cerf

Today's approach 1: Bolt-on protection



- Keep networks unchanged, deploy protections elsewhere
 - E.g., middleboxes, end host software, ..
- **Advantage: Immediately practical** 🤔
- **Disadvantage: Not a fundamental solution** 😞

Today's approach 2: Clean slate design



Internet re-imagined!



- **Redesign** the Internet to be secure
 - i.e., Future Internet Architectures
- **Advantage: Addresses the root cause**
- **Disadvantage: Expensive changes**



The best of both worlds?

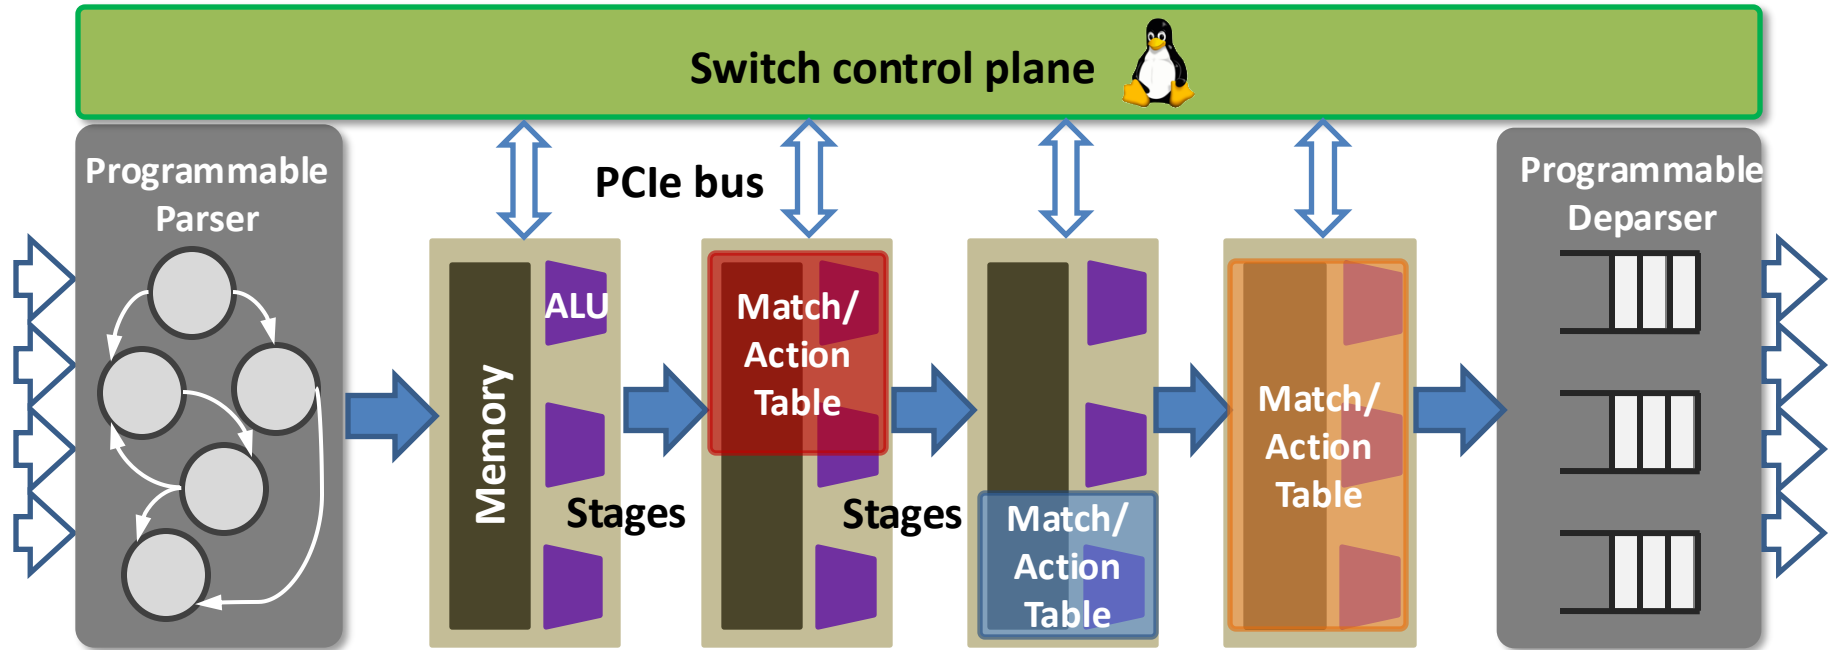
Built-in security

Compatibility



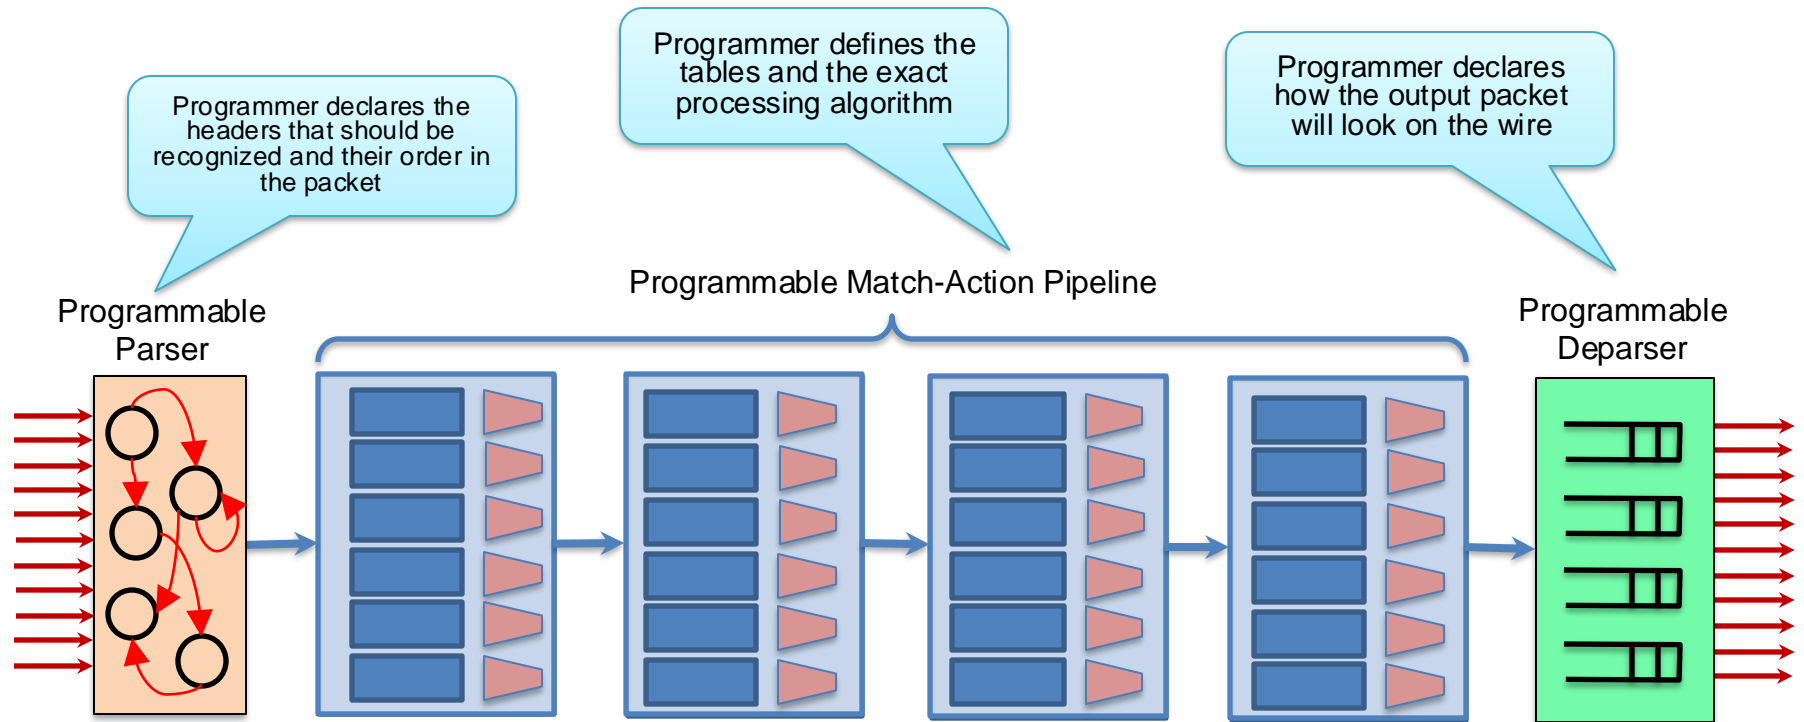
- Programmable In-network Security (Poise)
 - Architect security back to the network core
 - Without making intrusive modifications

Opportunity: Network programmability



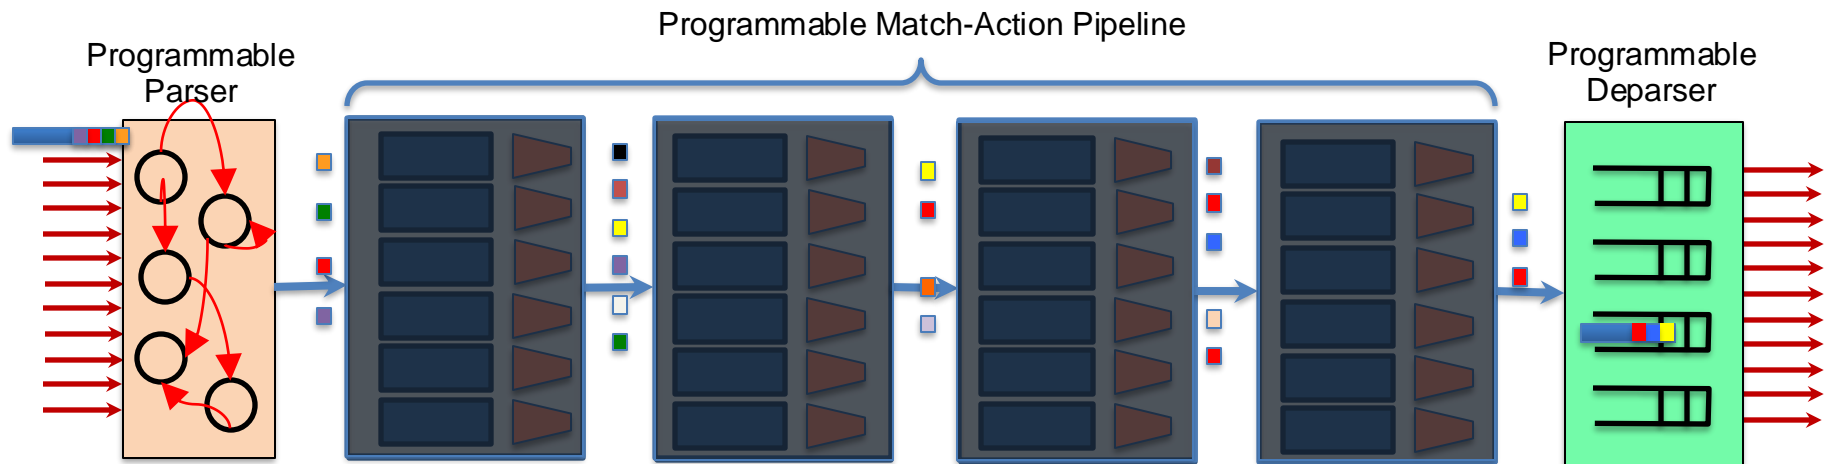
- Programmable switches
 - Controlled in high-level languages (e.g., P4)
 - Customized protocols, persistent state, custom match/actions
 - All at linespeed (Tbps)!
- We can develop defenses directly in the network!

Programming Protocol-independent Packet Processors



P4 in Action

- Packet is parsed into individual headers (parsed representation)
- Headers and intermediate results can be used for matching and actions
- Headers can be modified, added or removed
- Packet is deparsed (serialized)



V1Model Standard Metadata

```
struct standard_metadata_t {  
    bit<9>  ingress_port;  
    bit<9>  egress_spec;  
    bit<9>  egress_port;  
    bit<32> clone_spec;  
    bit<32> instance_type;  
    bit<1>  drop;  
    bit<16> recirculate_port;  
    bit<32> packet_length;  
    bit<32> enq_timestamp;  
    bit<19> enq_qdepth;  
    bit<32> deq_timedelta;  
    bit<19> deq_qdepth;  
    bit<48> ingress_global_timestamp;  
    bit<32> lf_field_list;  
    bit<16> mcast_grp;  
    bit<1>  resubmit_flag;  
    bit<16> egress_rid;  
    bit<1>  checksum_error;  
}
```

- **ingress_port** - the port on which the packet arrived
- **egress_spec** - the port to which the packet should be sent to

P4₁₆ Program Template (V1Model)

```
#include <core.p4>
#include <v1model.p4>

/* HEADERS */
struct metadata { ... }
struct headers {
    ethernet_t    ethernet;
    ipv4_t        ipv4;
}

/* PARSER */
parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t smeta) {
    ...
}

/* CHECKSUM VERIFICATION */
control MyVerifyChecksum(in headers hdr,
                        inout metadata meta) {
    ...
}

/* INGRESS PROCESSING */
control MyIngress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t std_meta) {
    ...
}
```

```
/* EGRESS PROCESSING */
control MyEgress(inout headers hdr,
                inout metadata meta,
                inout standard_metadata_t std_meta) {
    ...
}

/* CHECKSUM UPDATE */
control MyComputeChecksum(inout headers hdr,
                        inout metadata meta) {
    ...
}

/* DEPARSER */
control MyDeparser(inout headers hdr,
                  inout metadata meta) {
    ...
}

/* SWITCH */
V1Switch(
    MyParser(),
    MyVerifyChecksum(),
    MyIngress(),
    MyEgress(),
    MyComputeChecksum(),
    MyDeparser()
) main;
```

P4₁₆ Hello World (V1Model)

```
#include <core.p4>
#include <v1model.p4>
struct metadata {}
struct headers {}

parser MyParser(packet_in packet,
  out headers hdr,
  inout metadata meta,
  inout standard_metadata_t standard_metadata) {

  state start { transition accept; }
}

control MyVerifyChecksum(inout headers hdr, inout metadata
meta) {  apply { } }

control MyIngress(inout headers hdr,
  inout metadata meta,
  inout standard_metadata_t standard_metadata) {
  apply {
    if (standard_metadata.ingress_port == 1) {
      standard_metadata.egress_spec = 2;
    } else if (standard_metadata.ingress_port == 2) {
      standard_metadata.egress_spec = 1;
    }
  }
}
```

```
control MyEgress(inout headers hdr,
  inout metadata meta,
  inout standard_metadata_t standard_metadata) {
  apply { }
}

control MyComputeChecksum(inout headers hdr, inout metadata
meta) {
  apply { }
}

control MyDeparser(packet_out packet, in headers hdr) {
  apply { }
}

V1Switch(
  MyParser(),
  MyVerifyChecksum(),
  MyIngress(),
  MyEgress(),
  MyComputeChecksum(),
  MyDeparser()
) main;
```


P4₁₆ Hello World (V1Model)

```
#include <core.p4>
#include <v1model.p4>
struct metadata {}
struct headers {}

parser MyParser(packet_in packet, out headers hdr,
  inout metadata meta,
  inout standard_metadata_t standard_metadata) {
  state start { transition accept; }
}

control MyIngress(inout headers hdr, inout metadata meta,
  inout standard_metadata_t standard_metadata) {
  action set_egress_spec(bit<9> port) {
    standard_metadata.egress_spec = port;
  }
  table forward {
    key = { standard_metadata.ingress_port: exact; }
    actions = {
      set_egress_spec;
      NoAction;
    }
    size = 1024;
    default_action = NoAction();
  }
  apply { forward.apply(); }
```

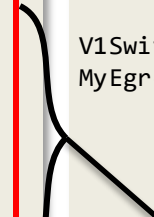
```
control MyEgress(inout headers hdr,
  inout metadata meta,
  inout standard_metadata_t standard_metadata) {
  apply { }
}

control MyVerifyChecksum(inout headers hdr, inout metadata
meta) { apply { } }

control MyComputeChecksum(inout headers hdr, inout metadata
meta) { apply { } }

control MyDeparser(packet_out packet, in headers hdr) {
  apply { }
}

V1Switch( MyParser(), MyVerifyChecksum(), MyIngress(),
MyEgress(), MyComputeChecksum(), MyDeparser() ) main;
```



Key	Action ID	Action Data
1	set_egress_spec ID	2
2	set_egress_spec ID	1

Poise vs. today's software defenses

Three key advantages

- Per-packet visibility
 - Can detect needle-in-a-haystack attacks
- Per-packet dynamicity
 - Can respond as fast as dynamic changing attacks
- Scale-free defense
 - Naturally scales with network size and speed


- At the same time, compatible with legacy networking

Roadmap: The switch as a defense platform



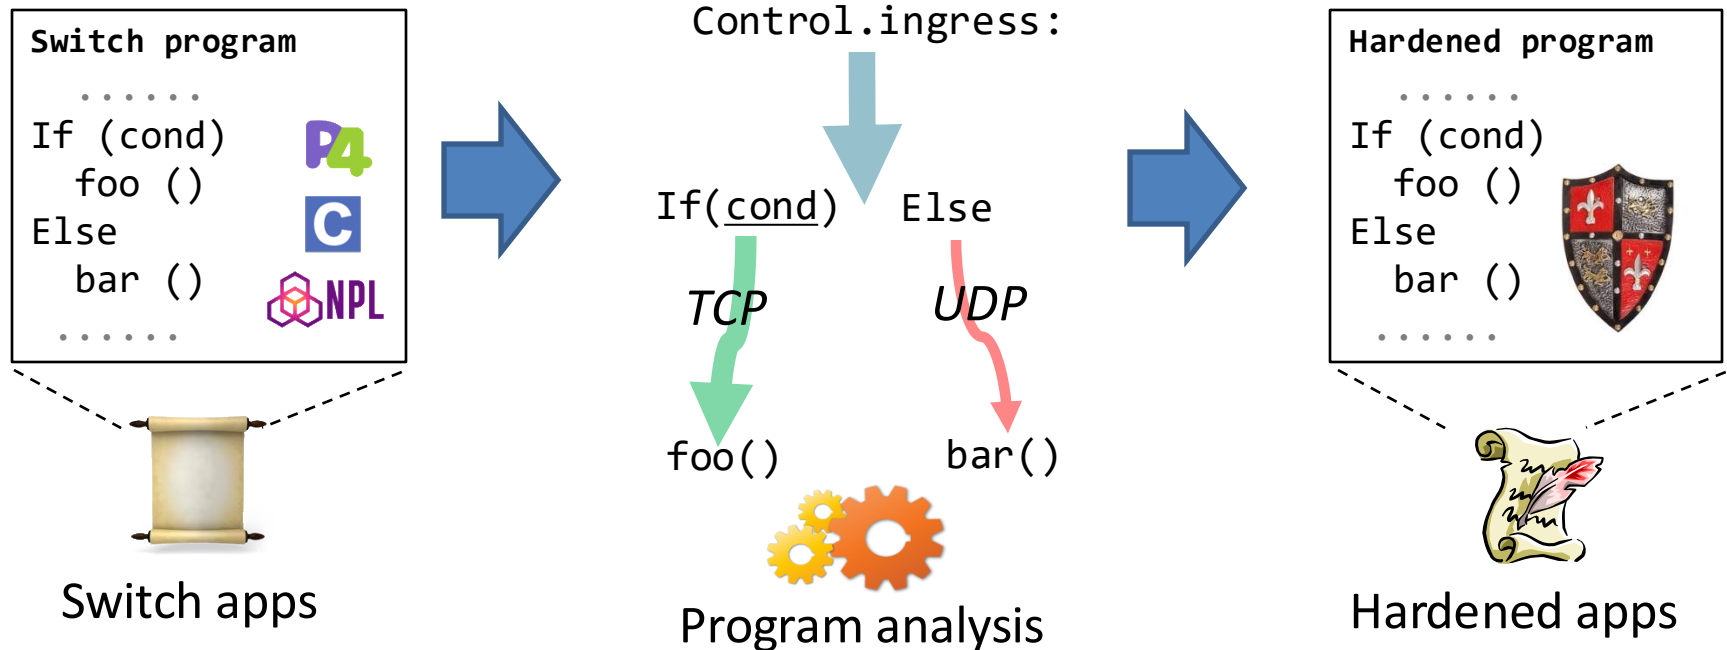
- Step 1: Design a range of switch-based security defenses
 - Switches dynamically activate the needed defenses
 - Progress so far:
 - Data exfiltration via covert channels (USENIX Sec'20)
 - Access control, RDMA security (USENIX Sec'20+'22)..

Roadmap: The network as a defense fleet



- Architect a wide range of defenses into network paths
 - Networks mitigate attacks as they route traffic
 - Progress so far: [Link flooding attacks \(USENIX Sec'21\)](#)

Roadmap: Securing the defenses

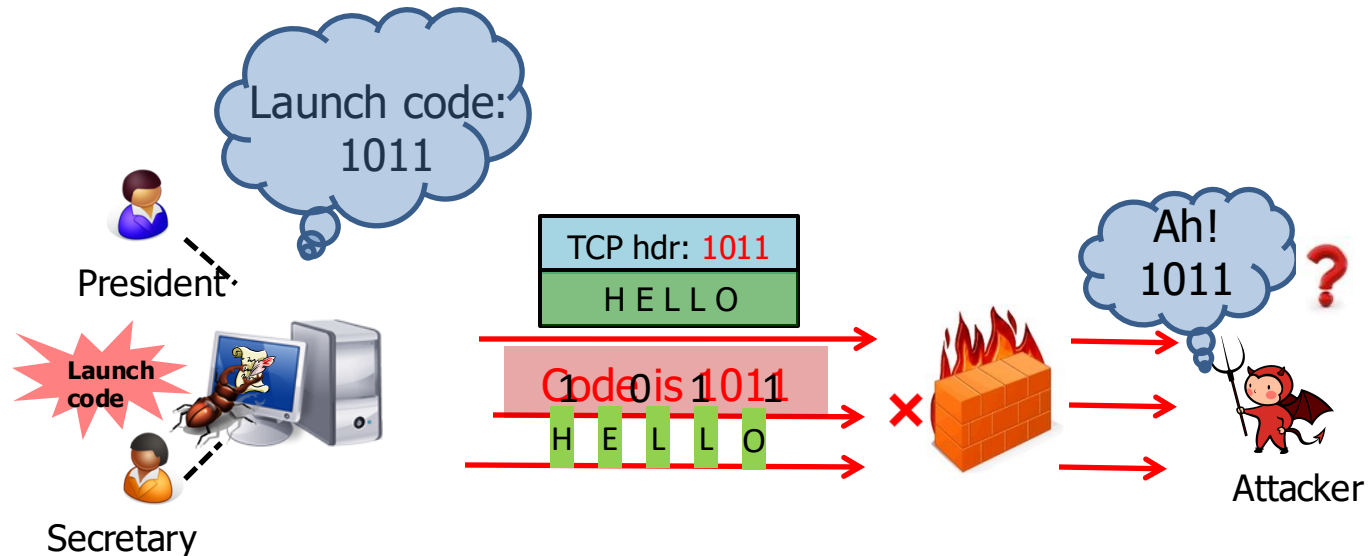


- Ensure that security apps are themselves secure
 - Program analyses to gain high assurance
 - Progress so far: [Adversarial trace discovery \(ASPLOS'21\)](#)

Outline


- ✓ • Programmable In-network Security
 - ➔ • The switch as a defense platform
 - Securing the switch apps
 - The network as a defense fleet
 - Future work
 - Summary


Motivation: Mitigating network covert channels



- Network covert channels:
 - Storage channels: changing the packet **header fields**.
 - E.g., TCP ISN (1997), TTL (2004), Partial ACK (2009)
 - Timing channels: changing the **timing** of packets
 - E.g., IP-layer (2004), TCP-layer (2008), PHY-layer (2014)

Detection algorithms are never perfect



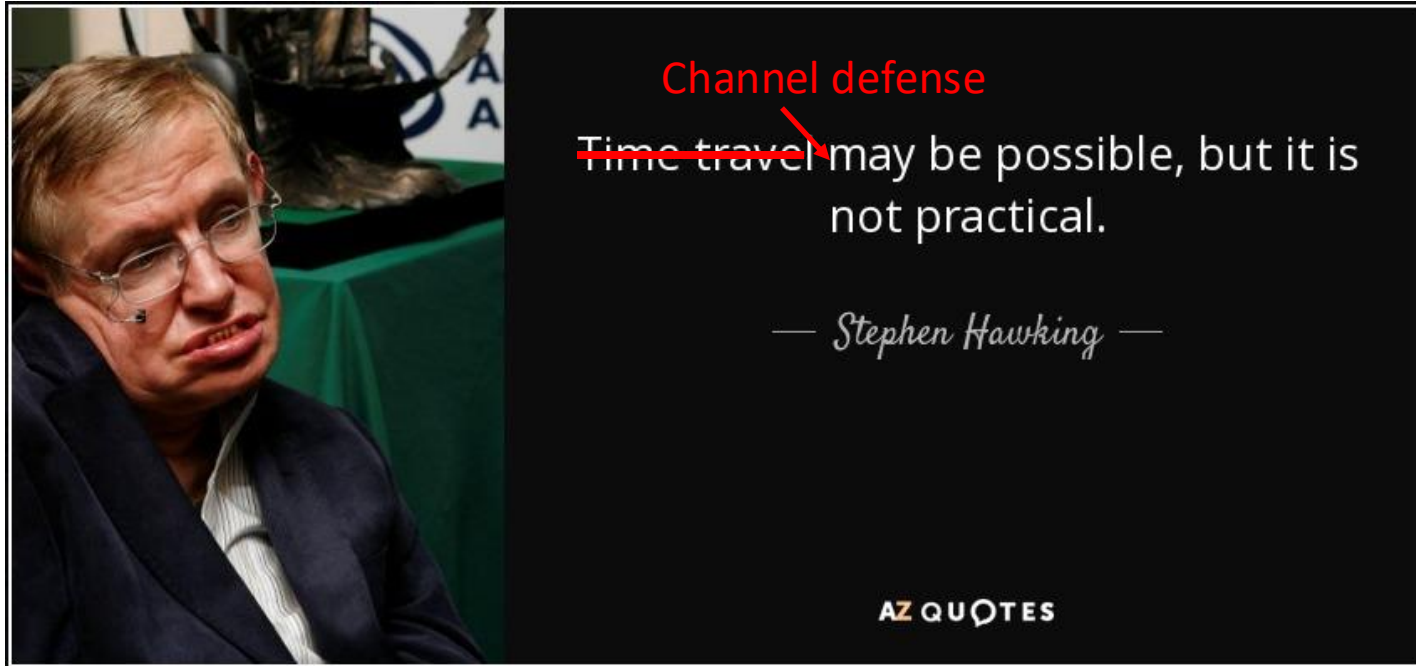
- Consider network timing channels 
- Packet timing is non-deterministic
- E.g., due to congestion, server processing speeds, ..
- Approach: statistical tests against known-good traces
 - Can result in false positives/negatives

Mitigation algorithms hurt performance




- Timing channel mitigation
 - E.g., Randomly delay packets to destroy timing modulation
- Storage channel mitigation
 - E.g., scrub header fields by adding offset to TCP SEQs

Implication: Performance penalty



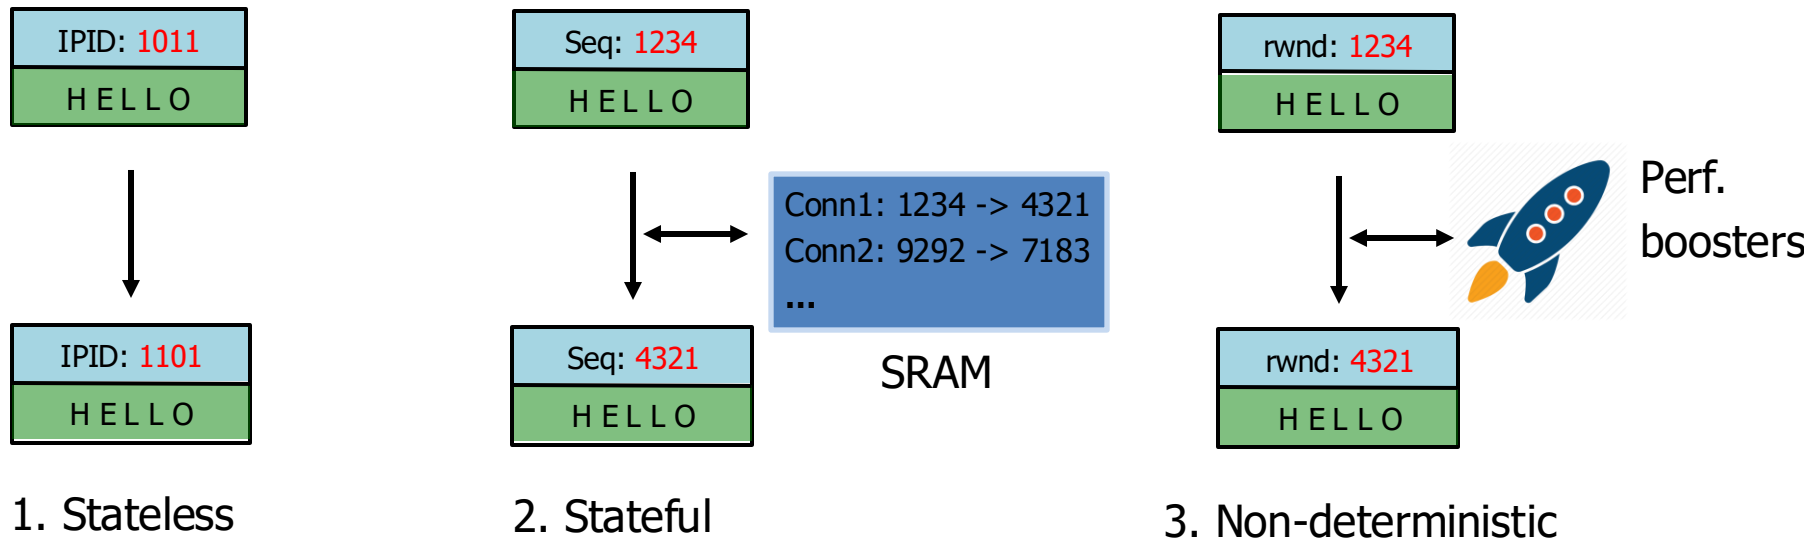
- Existing defenses incur heavy performance penalty
 - Needs to process every single packet
 - → Modern networks have high speeds (Tbps)
 - Imperfect detection leads to collateral damage
 - → Normal traffic may be adversely impacted

Goal: Performance-preserving defenses




- Per-packet scrubbing at hardware speeds
- Manipulate TCP congestion control to mask latency
- Batch operations in software for TCP reliability

Per-packet defense




- The P4 data plane performs per-packet defenses
 - Stateless: IPID, TTL, ...
 - Stateful: TCP sequence number, ACK, ..
 - Non-deterministic: receive window, timing channels..

Manipulating TCP congestion control



- Consider the timing channel defense:
 - Adding delay to packets disrupts TCP throughput
- **ACK boosters:**
 - Optimistically ACK packets from the switch
 - Create the illusion of the same RTT for the sender

ACK boosters



- Creates the illusion of the same RTT despite extra delay
 - (1) Measures (actual) RTT in data plane
 - (2) Generates ACK optimistically to trigger more data
 - (3) Caches extra packets for reliability
 - (4)-(5) Serves the client from the cache
- Similar as a “performance-enhancing TCP proxy”
 - But accelerated in P4 hardware

Hardware/software codesign

Switch control plane

- Computation: Batch operations
- Memory: Growing state


- Minimize crosstalk

- Computation: Per-packet operations
- Memory: Constant state

Switch data plane



Pkt cache



Stat. tests



RTT



Flow state




IPDs



Header ops

- Challenge: P4 has a restricted programming model
 - Cannot perform statistical tests, packet caching ..
- Solution: Division of labor between SW+HW
 - Principles are generally applicable to other defenses

Other techniques: Computing IPDs




- Challenges: for computing inter-packet gaps
 - Per-packet operations needed
 - But also accumulates state very fast (2B per packet)
- Solution: two approximation techniques

Evaluation setup


- NetWarden prototype:
 - Runs in Tofino Wedge 100BF-32X switch.
 - 2500 LoC of P4 + 3000 LoC of C+Python
- Threat model:
 - A compromised server + a trusted P4 switch running NetWarden
 - Leak a 2048-bit RSA key via covert channel.
- Real world applications:
 - Apache servers, FTP servers, Nodejs servers
- Baseline:
 - No defense: No covert channel defenses are deployed.
 - Naïve defense: Covert channel defenses without performance preservation

Defense effectiveness




- Naïve defense: renders decoding to a random guess.
- NetWarden: very close to a random guess.
- → Mitigates covert channels effectively.

Performance preservation



- Naïve defense incurs 25% performance penalty.
- NetWarden only has 1% performance deviation
- → Mitigates channels with minimal performance loss.


Other results



System
scalability



System
overhead



Different TCP
variants



Complex
Applications




Self-defense
techniques

Outline


- ✓ • Programmable In-network Security
 - ✓ • The switch as a defense platform
 - ➔ • The network as a defense fleet
 - Securing the defenses
- Future work
- Summary

Motivation: Network-wide attacks




- Individual switches have local views, but global views and actions are needed
- Example: **Link-flooding attacks [USENIX Security'21]**
 - Congest critical links to take down remote victim
 - Attackers can easily change traffic type, link target, ..

State of the art: SDN defenses



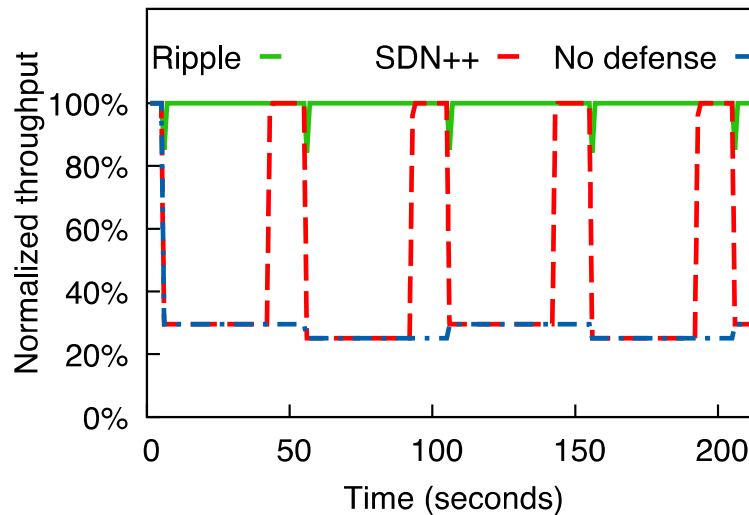
- A central controller samples traffic, computes decisions, and installs rerouting/filtering rules
- Limitation: Cannot handle dynamic attacks
 - Changing botnet composition
 - Changing traffic patterns
 - Changing target links

Ripple: A fully decentralized defense



- Programmable switches coordinate in data plane
- A panoramic view as defined in a policy language
- Switches extract local signals, propagate them across the network, and reconstruct the global view

Ripple mitigates fast-changing attacks




- Setup: Dynamic attacks based on “Crossfire” [SP ‘13]
- Baseline solutions: SDN variants
- Ripple can mitigate attacks in real time, whereas SDN defenses constantly lag behind.

Outline

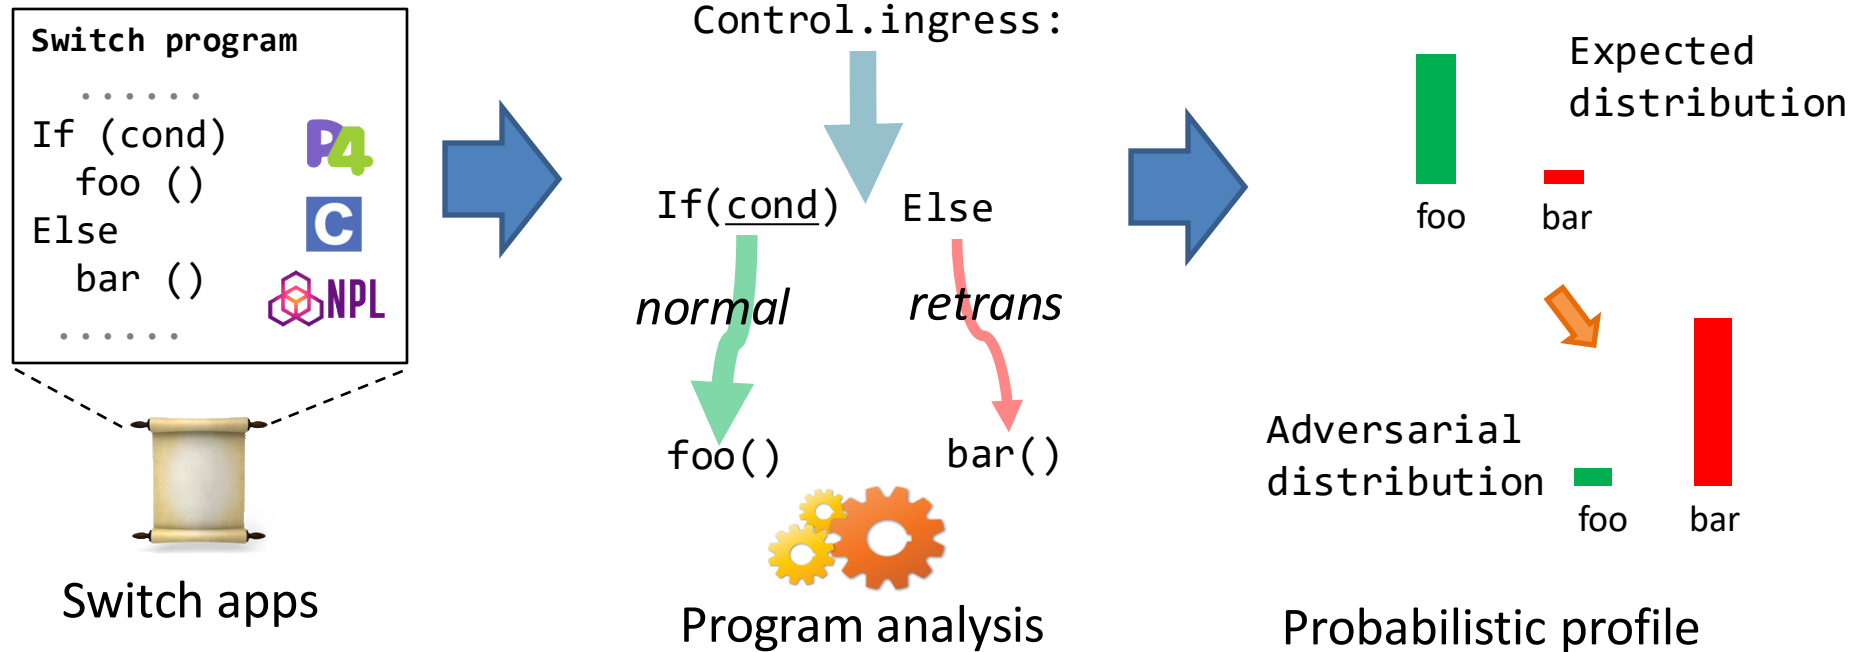
- ✓ • Programmable In-network Security
 - ✓ • The switch as a defense platform
 - ✓ • The network as a defense fleet
 - ➔ • Securing the switch apps
- Other work
- Summary

Motivation: Adversarial testing



- Automatically identify “adversarial traces” for a P4 system?
 - Example: We’ve manually done this for NetWarden
 - Adversarial traces exist for other P4 systems, too

Worst-case behavior analysis?



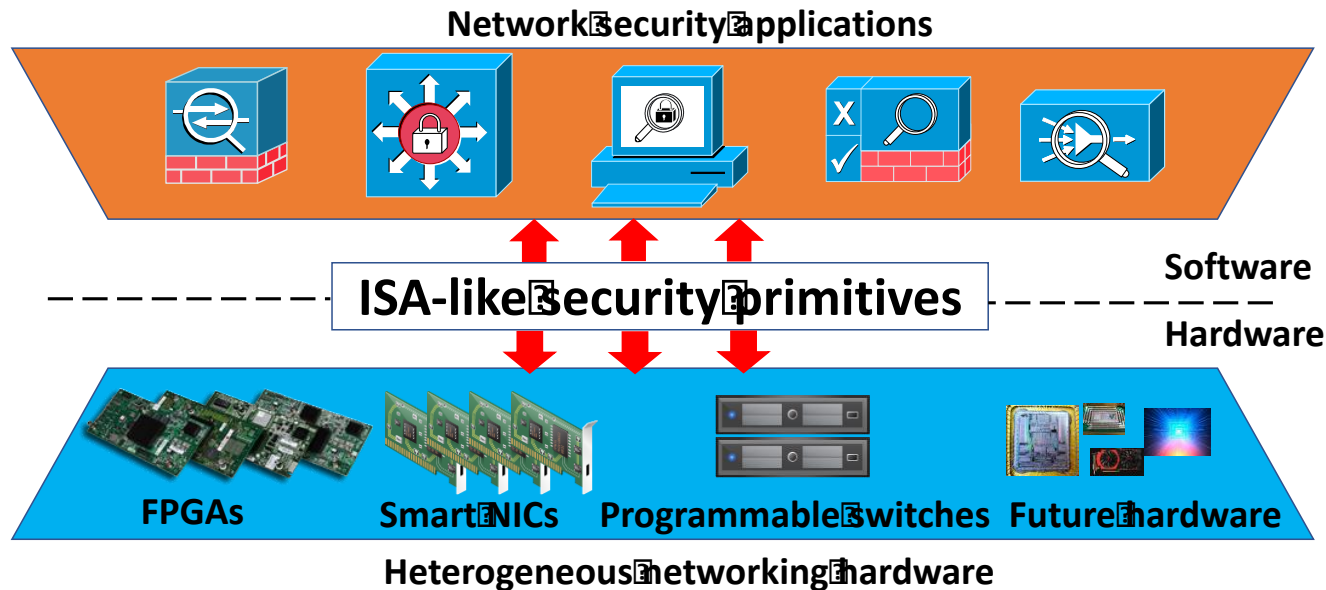
- Intuition: Network programs have complex behaviors, and some are easier to process than others.

Where to go from here?




- Programmable In-network Security (Poise)
 - The switch as a defense platform
 - Securing the switch apps
 - The network as a defense fleet

The Poise “network security stack”




- The “narrow waist” of the network security stack
 - ISA-like primitives for security
 - Small building blocks, verifiable
 - Recomposable for different tasks
 - Dynamically swapped in and out for defense

Need to add runtime flexibility!




- Operators compile and reflash a new program
 - Causes data plane disruption, so need to provision for that
 - Drain and undrain traffic for any data plane changes
 - Expensive to plan and perform, constrains change velocity

Runtime programmable networks



- Runtime network (re)programming end-to-end
- No downtime, zero packet loss, consistency guarantees
- Users can inject real-time “network extensions”

Why runtime programmable networks?



- Use case: Real-time security defense
 - Network devices swaps defenses in and out
 - Defenses dynamically migrate across the network
 - They shapeshift in real time for changing attacks

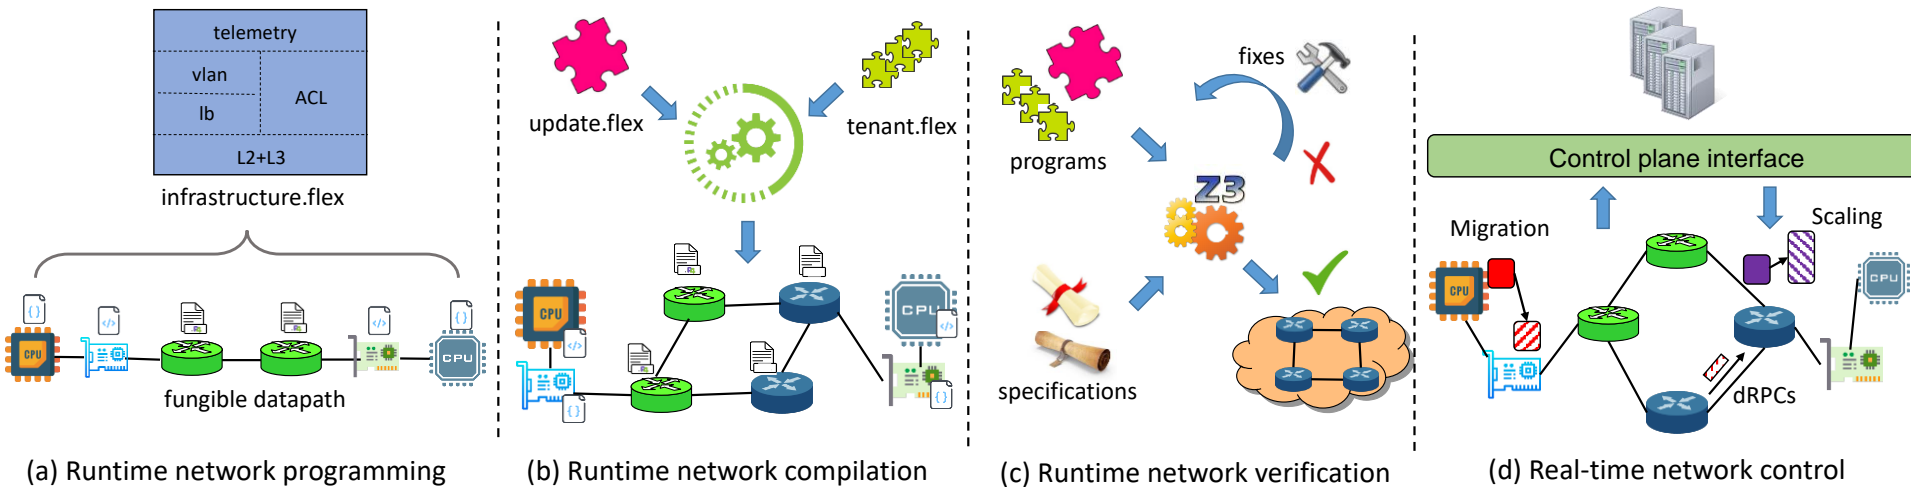
Why runtime programmable networks?

- Just-in-time network optimization
 - Common mode: Basic routing, low latency
 - JIT optimization based on applications and workloads
- Tenant-specific extensions
 - Tenant directly customize network logic
- Incremental infrastructure upgrades
 - Coordinated changes at NICs, switches, and hosts
 - Ex: new congestion control algorithms

Why now?

- Individual targets are becoming runtime programmable
 - Ex: Runtime programmable switches (NSDI'22)
 - Dynamically add/remove/modify tables, control flow, parsers
 - Prototyped on a 12.8Tbps ASIC
 - Ex: Runtime programmable SmartNICs (ongoing)
 - Ex: Host networking based on eBPF

The overall vision



- Brand new challenges and opportunities ahead!

Summary

- Motivation: Next-generation network security
- Approach: Programmable In-network security
 - Architect security back into the network foundation
- We have already made some progress:
 - A switch as a defense platform
 - A network as a defense fleet
 - Securing the defenses
- More opportunities await!

Thank you!